

Report 2: DBMS - Postgresql, Cockroach, Yugabyte - Basic

7.

Název tématu	Měření výkonu DB enginů
Určení	PVS
Anotace	<i>Realizace sady experimentů pro porovnání výkonu DB s implementovaným rozhraním typu Postgres (Postgres, Yagabyte, Cocroach).</i>
Výstupy práce	<i>Výzkumná zpráva + software</i>
Požadavky na studenta	<i>Student alespoň 3. ročníku KB, případně student s nadprůměrnými znalostmi jazyku Python</i>
Navrhovatel	<i>Prof. Dr. Ing. Alexandr ŠTEFEK</i>

How I understand what I have to do:

Topic Title:	Measuring the Performance of DB Engines
Application:	PVS
Annotation	Implementation of a set of experiments to compare the performance of DB with the implemented interface of the Postgres type (Postgres, Yagabyte, Cocroach).
Outputs	Research Report + Software
Student Requirements	Student of at least the 3rd year of KB, or a student with above-average knowledge of Python
Applicant:	Prof. Dr. Ing. Alexandr ŠTEFEK

Good afternoon,

Firstly, slow down with uois, there will be new version you can experiment with soon. I expect next week.
Secondly, try to read documentation about cocroach and yugabyte regarding its installation in Docker environment.

Alexandr Stefek

Dobrý den,

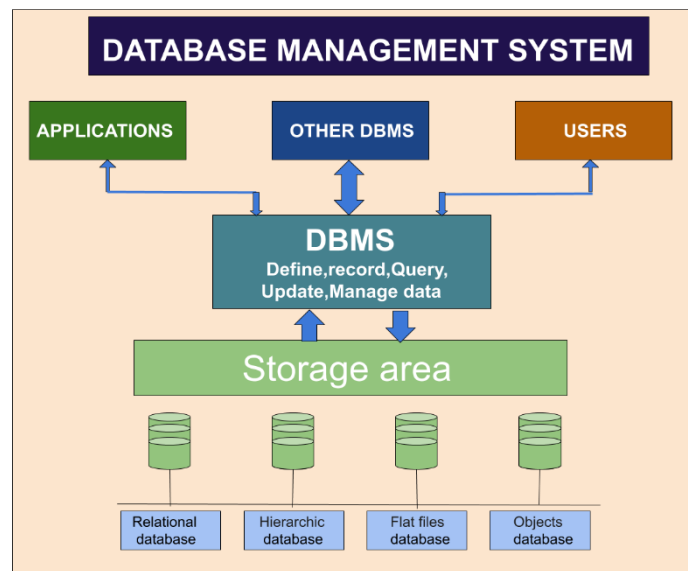
1. Work with docker-compose.yaml configuration, this allows you create a net with services on it capable to see each other out of box.
2. Try check this: <https://github.com/dbist/cockroach-docker/blob/main/cockroach-prometheus/docker-compose.yml>

Hope this helps
Alexandr Štefek

1. Definitions

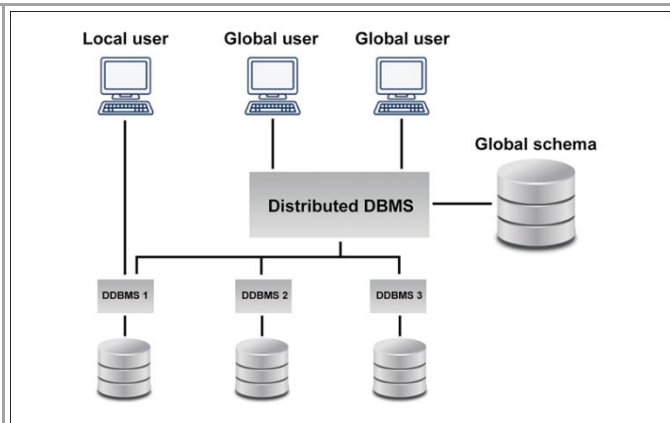
What is DBMS?

Database Management Systems (DBMS) are software systems used to store, retrieve, and run queries on data. A DBMS serves as an interface between an end-user and a database, allowing users to create, read, update, and delete data in the database.



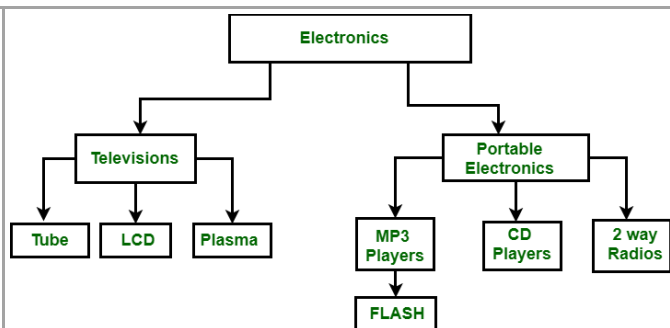
Distributed database management system

A distributed DBMS is a set of logically interrelated databases distributed over a network that is managed by a centralized database application. This type of DBMS synchronizes data periodically and ensures that any change to data is universally updated in the database.




Hierarchical database management system

Hierarchical databases organize model data in a tree-like structure. Data storage is either a top-down or bottom-up format and is represented using a parent-child relationship.



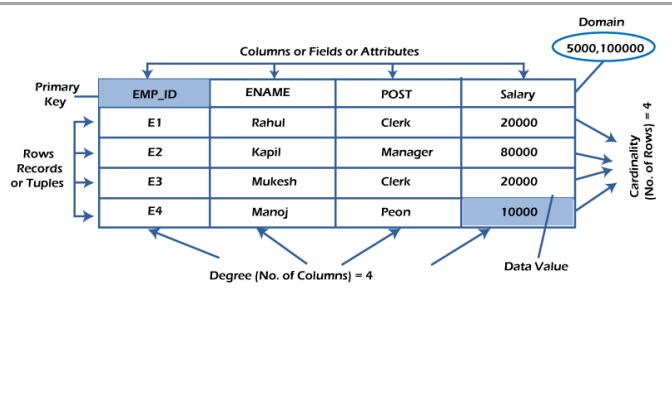
Network database management system

The network database model addresses the need for more complex relationships by allowing each child to have multiple parents. Entities are organized in a graph that can be accessed through several paths.



Relational database management system

Relational database management systems (RDBMS) are the most popular data model because of its user-friendly interface. It is based on normalizing data in the rows and columns of the tables. This is a viable option when you need a data storage system that is scalable, flexible, and able to manage lots of information.



1.2 DB-Engines

A database engine is a component of software that facilitates working with databases. It controls how information in a database is saved, retrieved, and altered. Database engines act as a bridge between the database's data and the user, facilitating interaction with the information stored there

A database engine might handle multiple users, transactions, throughput, **As I mentioned in previous report**, buffers and caches, ACID (atomicity, consistency, isolation, durability), as well as different isolation levels.

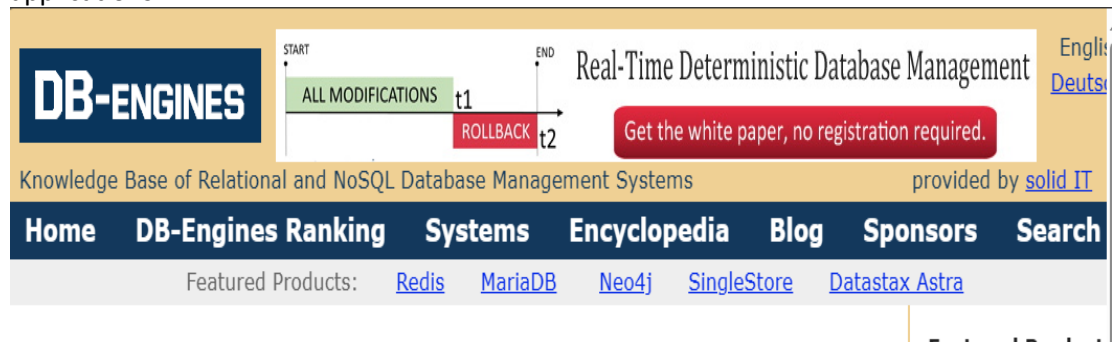
- A read may pull data from memory, remote databases, and multiple tables on disk processing it using SQL through multiple explicit and/or implicit code paths in order to present it to the requesting application.
- A create may allocate storage, provision structures, assign values, and do its own processing before storing data. Etc.

DB-Engines website

DB-Engines is an initiative to collect and present information on database management systems (DBMS)..

The DB-Engines Ranking is a list of DBMS ranked by their current popularity. The list is updated monthly.

DB-Engines has been created and is maintained by solid IT - an Austrian IT consulting company with a special focus on software development, consulting and training for database-centric applications.



Method of calculating the scores of the DB-Engines Ranking

DB-Engines Ranking - Trend of CockroachDB vs. PostgreSQL vs. YugabyteDB Popularity
[CockroachDB vs. PostgreSQL vs. YugabyteDB Comparison \(db-engines.com\)](#)

Relational Database Engines

Relational database engines are the most widely used type of database engines.

Relational database engines store data in tables, with each table representing a different entity or object. Each table has a unique primary key that is used to identify the rows in the table. Tables can be related to each other using foreign keys, which are used to establish relationships between tables.

To retrieve data from a relational database, users use SQL queries. SQL is a standard language used to manipulate relational databases. SQL queries are used to select, insert, update, and delete data from the database.

Performance.

Relational databases perform well with intensive read/write operations on small to medium datasets. They also offer improved speed of data retrieval by adding indexes to data fields to query and join tables. However, the performance may suffer when the amount of data and user requests grows.

2. Ideas from case studies:

Mainly, studies use the same criteria to determine performance: the execution time of the different operations

The operations choosed in the studies are **INSERT**, **DELETE**, **UPDATE**, and **SELECT**. These four operations are the primary operations of the DBMSs. Therefore, the performance of different DBMSs is based on these operations.

While another research studies the performance in terms of throughput and latency. The performance is measured by measuring throughput and latency for four queries.

- **S1:** testing these operations with **increasing number of records**. This study is based on 1, 100, 500, 1 000, 5 000, 10 000, 25 000, and 50 000 records. The study shows that the execution time differs in databases when executing with different numbers of records.

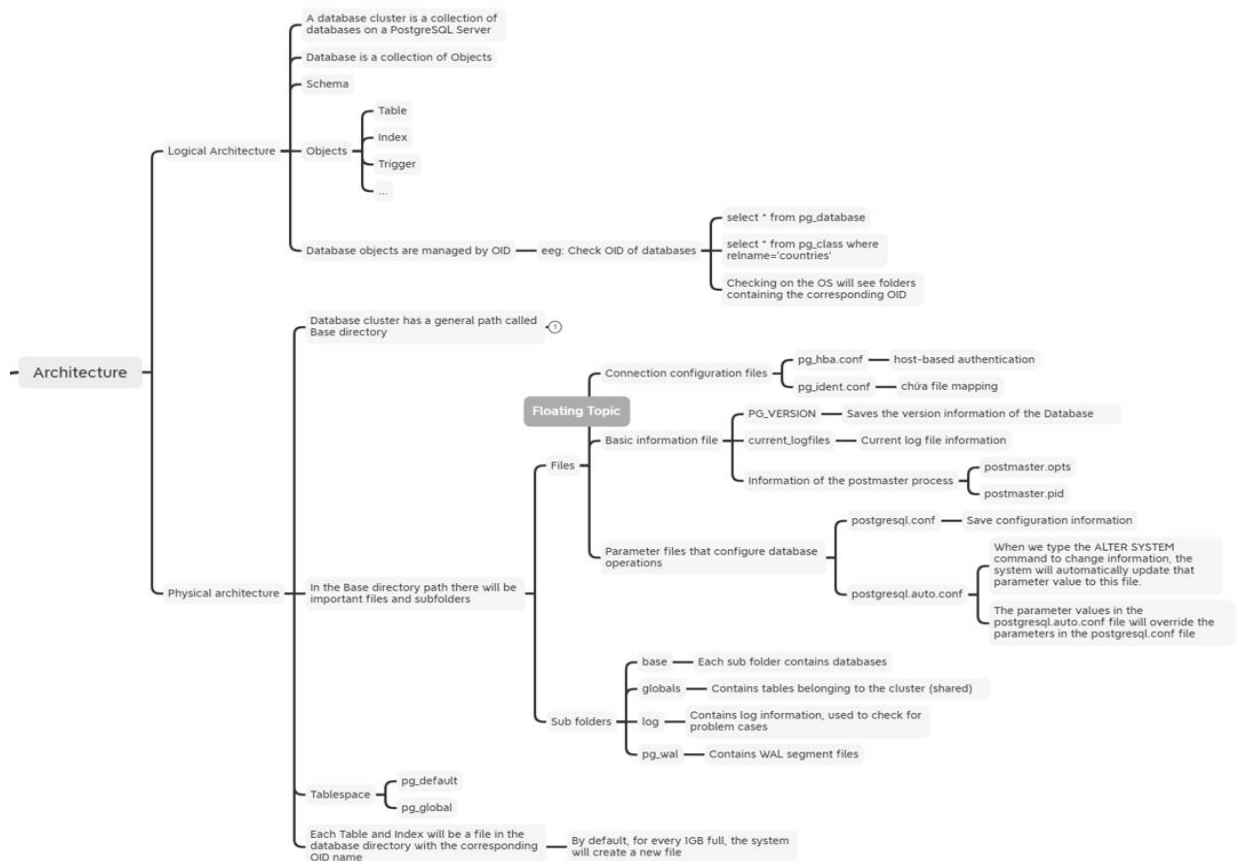
- **S2:** the number of records varies from 1 000 to 100 000. The study examined the databases by conducting tests on the CRUD operations, creating, reading, updating, and deleting. **Each test has been conducted five times on the same operation and the same number of records.** Finally, the result is entirely based on the **average execution times of the tests.** Furthermore, the operations were conducted to change and update this information.
- **S3:** Examines the performance by conducting tests on operations of **insertion and search.**
- **S4:** databases contain more than 100 000 records The execution time of these 30 sets for the four operations has been presented. The differences are slight in terms of the execution time for these sets. Therefore the average execution time for each operation has been presented. This execution average time shows the differences between the operations.

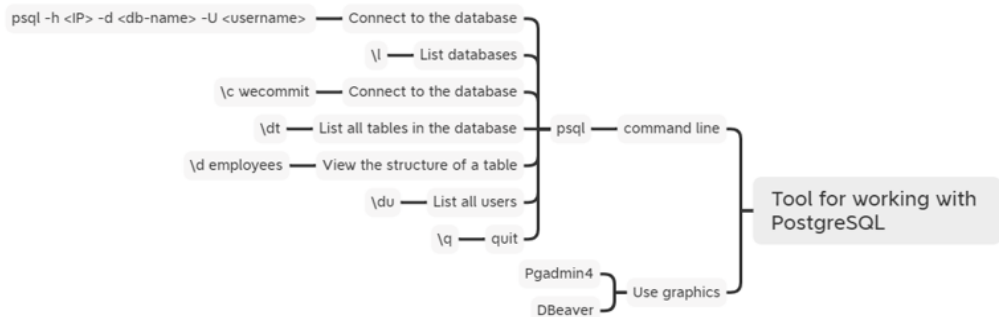
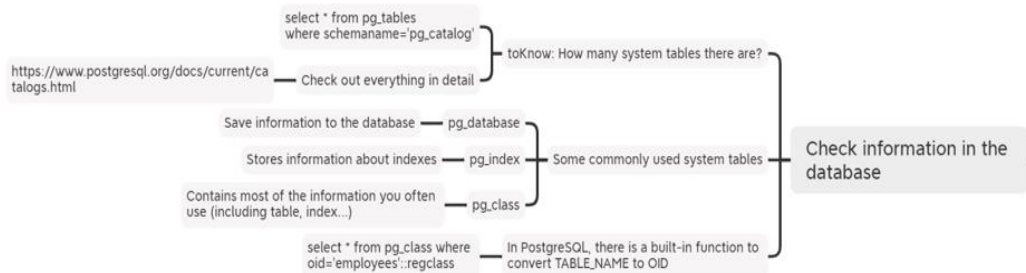
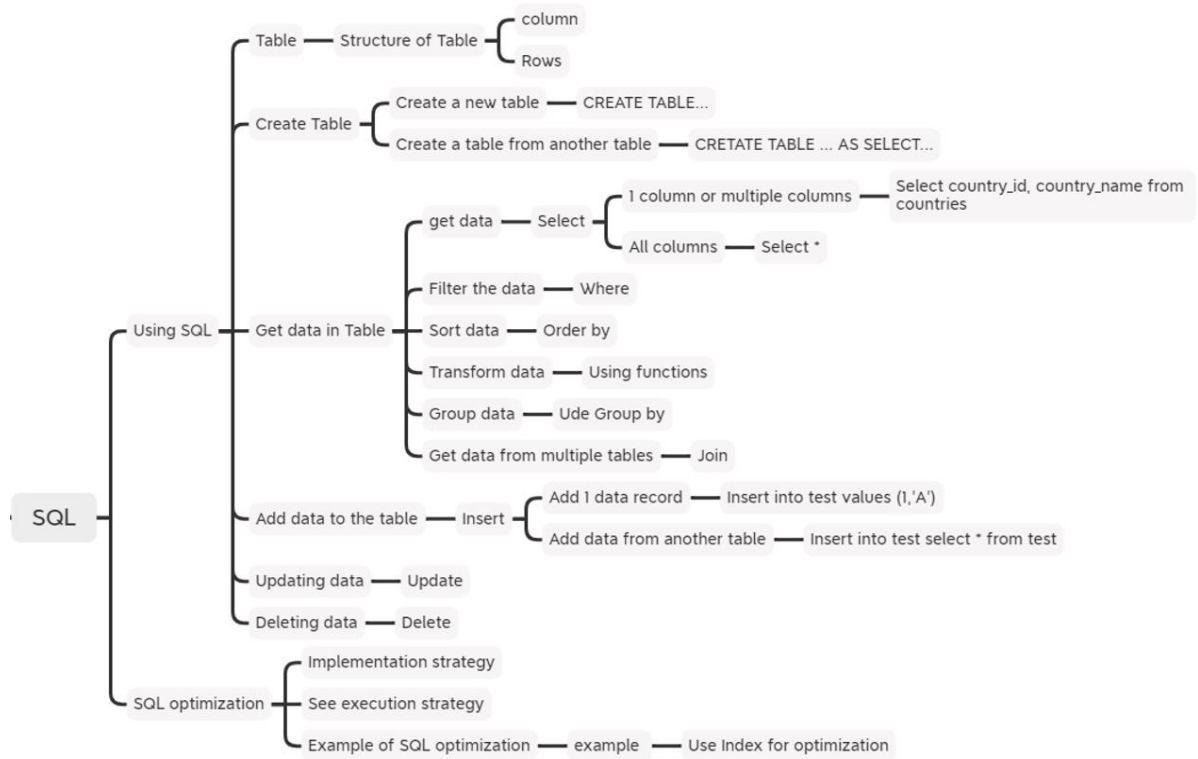
3. Jump in Postgresql, Cockroach, Yugabyte - Basic Introduce + Installation- preparing for the tests

3.1 Postgresql

PostgreSQL is an advanced, enterprise-class, and open-source relational database system. PostgreSQL supports both SQL (relational) and JSON (non-relational) querying. PostgreSQL is a highly stable database backed by more than 20 years of development by the open-source community.

Postgresql mindmap





Docker

Main reasons for using docker database image for your local development and why I think you should do it too:

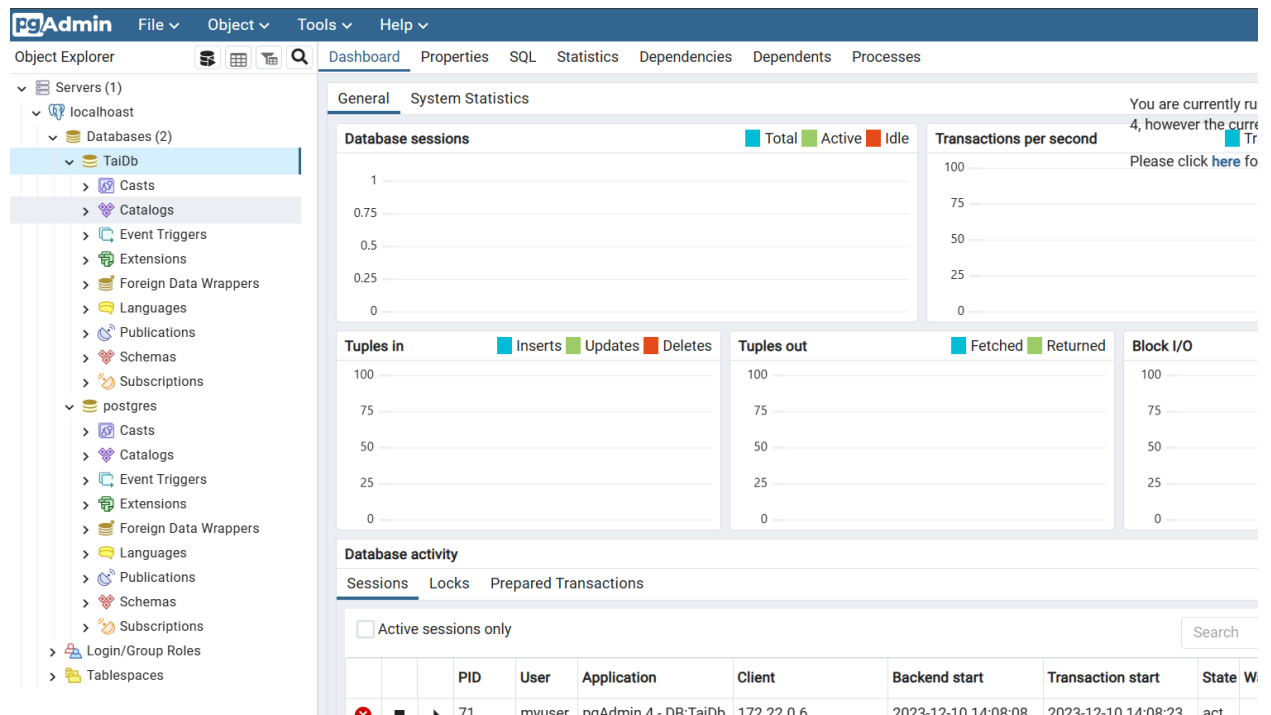
- It makes it easy for local development setup. A simple abstraction using a one-line setup for getting your database set up!
- It does not require you to have any databases installed at all (or a specific version of the database) on your computer.
- It isolates the database environment. It allows me to run multiple versions and instances at the same time without getting in trouble.

```
services:
##### postgres #####
postgres:
  container_name: container-pg
  image: "postgres:16.0"
  hostname: localhost
  ports:
    - "5432:5432"
  environment:
    POSTGRES_USER: myuser
    POSTGRES_PASSWORD: data
    POSTGRES_DB: TaiDb
  networks:
    - roachnet
  restart: unless-stopped

##### pgadmin #####
pgadmin:
  container_name: container-pgadmin
  image: "dpage/pgadmin4:7.8"
  environment:
    PGADMIN_DEFAULT_EMAIL: admin@admin.com
    PGADMIN_DEFAULT_PASSWORD: root
  depends_on:
    - postgres
  ports:
    - "5050:80"
  networks:
    - roachnet
  restart: unless-stopped
```

```
PS D:\Documents\Unob_7\STC\STC_code\Cockroach\Tai_cockroach> docker ps -a
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS                               NAMES
7fe398c591a5   dpage/pgadmin4:7.8                 "/entrypoint.sh"        5 minutes ago Up 5 minutes  443/tcp, 0.0.0.0:5050->80/tcp        container-pgadmin
546c4ab26159   postgres:16.0                      "docker-entrypoint.s..." 5 minutes ago Up 5 minutes  0.0.0.0:5432->5432/tcp              container-pg
3498c7b66cd    cockroachdb/cockroach:v23.1.11     "/cockroach/cockroac..." 41 hours ago  Up 26 minutes  8080/tcp, 0.0.0.0:8081->8081/tcp, 26257/tcp, 0.0.0.0:26258->26258/tcp  roach2
88ed1ad36f1    cockroachdb/cockroach:v23.1.11     "/cockroach/cockroac..." 41 hours ago  Up 26 minutes  0.0.0.0:8080->8080/tcp, 0.0.0.0:26257->26257/tcp  roach1
b01ac6864769   cockroachdb/cockroach:v23.1.11     "/cockroach/cockroac..." 41 hours ago  Up 26 minutes  8080/tcp, 0.0.0.0:8082->8082/tcp, 26257/tcp, 0.0.0.0:26259->26259/tcp  roach3
cf38318369e    hello-world                         "/hello"                2 weeks ago   Exited (0) 2 weeks ago              infallible_archimedes

PS D:\Documents\Unob_7\STC\STC_code\Cockroach\Tai_cockroach>
```



3.2 Cockroach

CockroachDB is a distributed SQL database built on a transactional and strongly-consistent key-value store. It scales horizontally; you can quickly and simply add nodes to your cluster to handle more traffic.

A wire protocol is the format for interactions between a database server and its clients. It encompasses authentication, sending queries, receiving responses, and so on. It is a description of the exact bytes sent and received by servers and clients. It does NOT encompass the actual query language itself, let alone database semantics.

CockroachDB is wire-compatible with PostgreSQL clients, meaning that if you have a PostgreSQL client library for your favorite programming language, you can use that to connect to CockroachDB.

The database features and SQL dialect itself are not 100% compatible with PostgreSQL, and likely never will be.

Also, even if all the SQL you do want to use is supported, it is likely a sub-optimal idea to just port some SQL over to the CockroachDB platform without additional investigation. Some issues, like interleaving values from different tables, should be examined to ensure good performance on CockroachDB.

And even setting aside issues like that, the access patterns you'd use for a multi-region distributed database are going to be a bit different than a high-availability PostgreSQL cluster. Data locality and all that.

The wide availability of client drivers just means you can get started quickly on popular platforms.

CockroachDB architecture terms

Cluster

- A group of interconnected CockroachDB nodes that function as a single distributed SQL database server. Nodes collaboratively organize transactions, and rebalance workload and data storage to optimize performance and fault-tolerance.
- Each cluster has its own authorization hierarchy, meaning that users and roles must be defined on that specific cluster.
- A CockroachDB cluster can be run in CockroachDB Cloud, within a customer Organization, or can be self-hosted.

Node: An individual instance of CockroachDB. One or more nodes form a cluster.

Range: CockroachDB stores all user data (tables, indexes, etc.) and almost all system data in a sorted map of key-value pairs. This keyspace is divided into contiguous chunks called ranges, such that every key is found in one range.

Replica: A copy of a range stored on a node. By default, there are three replicas of each range on different nodes.

Leaseholder: The replica that holds the "range lease." This replica receives and coordinates all read and write requests for the range. For most types of tables and queries, the leaseholder is the only replica that can serve consistent reads (reads that return "the latest" data).

Raft protocol: The consensus protocol employed in CockroachDB that ensures that your data is safely stored on multiple nodes and that those nodes agree on the current state even if some of them are temporarily disconnected.

Raft leader: For each range, the replica that is the "leader" for write requests. The leader uses the Raft protocol to ensure that a majority of replicas (the leader and enough followers) agree, based on their Raft logs, before committing the write. The Raft leader is almost always the same replica as the leaseholder.

Raft log: A time-ordered log of writes to a range that its replicas have agreed on. This log exists on-disk with each replica and is the range's source of truth for consistent replication.

Docker file

```
##### cockroach #####
roach1:
  image: cockroachdb/cockroach:v23.1.11
  container_name: roach1
  hostname: roach1
  restart: unless-stopped

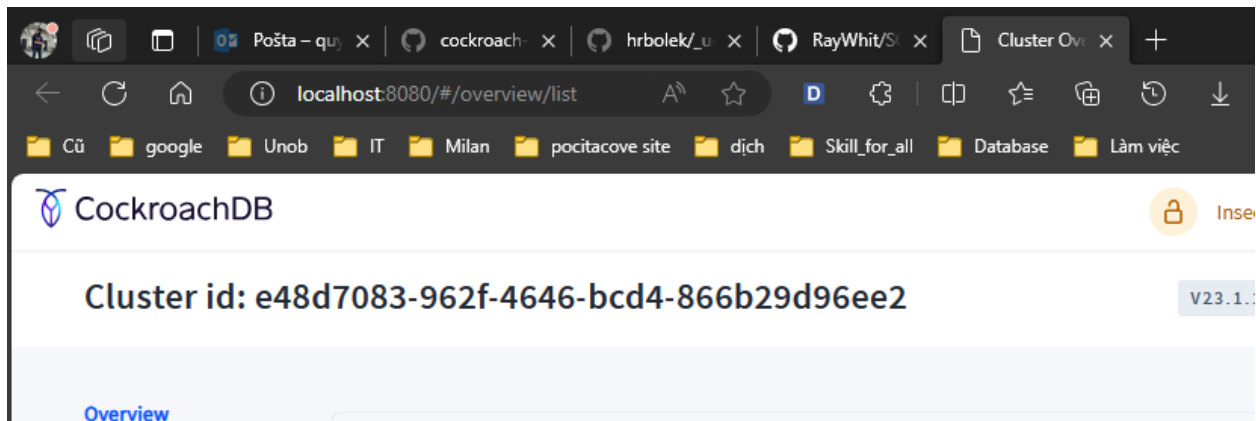
  # joins node with the others creating a cluster
  command: start --advertise-addr=roach1:26357 --http-addr=roach1:8080 --listen-addr=roach1:26357
  # command: start --advertise-addr=roach1 --http-addr=roach1 --listen-addr=roach1

  networks:
    - roachnet

  ports:
    - "26257:26257"
    - "8080:8080"



  volumes:
    - roach1:/cockroach/cockroach-data
```




```
Windows PowerShell
PS C:\Users\dangq> docker exec -it roach1 ./cockroach --host=roach1:26357 init --insecure
Cluster successfully initialized
PS C:\Users\dangq> |
```



Run the workload in 5mins and see the UI

```
PS C:\Users\dangq> cockroach workload init movr 'postgresql://root@localhost:26257?sslmode=disable';
I231208 21:26:08.358645 1 workload\cli\run.go:633 [-] 1 random seed: 10242531485974583653
I231208 21:26:08.549766 1 ccl\workloadccl\fixture.go:315 [-] 2 starting import of 6 tables
I231208 21:26:11.608347 56 ccl\workloadccl\fixture.go:492 [-] 4 imported 413 B in user_promo_codes table (5 rows,
0 index entries, took 2.2980378s, 0.00 MiB/s)
I231208 21:26:11.795986 51 ccl\workloadccl\fixture.go:492 [-] 5 imported 4.8 KiB in users table (50 rows, 0 index
entries, took 2.4868133s, 0.00 MiB/s)
I231208 21:26:11.799245 52 ccl\workloadccl\fixture.go:492 [-] 6 imported 3.2 KiB in vehicles table (15 rows, 15 i
dex entries, took 2.490073s, 0.00 MiB/s)
I231208 21:26:11.862400 55 ccl\workloadccl\fixture.go:492 [-] 7 imported 218 KiB in promo_codes table (1000 rows,
0 index entries, took 2.5521735s, 0.08 MiB/s)
I231208 21:26:12.208138 53 ccl\workloadccl\fixture.go:492 [-] 8 imported 153 KiB in rides table (500 rows, 1000 i
dex entries, took 2.8983735s, 0.05 MiB/s)
I231208 21:26:12.210013 1 ccl\workloadccl\fixture.go:323 [-] 9 imported 451 KiB bytes in 6 tables (took 3.6579555
, 0.12 MiB/s)
I231208 21:26:12.348307 1 workload\workloadsql\workloadsql.go:148 [-] 10 starting 8 splits
I231208 21:26:12.956778 1 workload\workloadsql\workloadsql.go:148 [-] 11 starting 8 splits
I231208 21:26:14.026119 1 workload\workloadsql\workloadsql.go:148 [-] 12 starting 8 splits
PS C:\Users\dangq>
PS C:\Users\dangq> cockroach workload run movr --duration=5m 'postgresql://root@localhost:26257?sslmode=disable';
I231208 21:26:26.653461 1 workload\cli\run.go:633 [-] 1 random seed: 18390389253575686851
I231208 21:26:26.653461 1 workload\cli\run.go:432 [-] 2 creating load generator...
I231208 21:26:26.655467 1 workload\cli\run.go:471 [-] 3 creating load generator... done (took 2.0064ms)
__elapsed__errors__ops/sec(inst)__ops/sec(cum)__p50(ms)__p95(ms)__p99(ms)__pMax(ms)
1.0s 0 59.4 59.6 10.5 24.1 26.2 83.9 readVehicles
1.0s 0 1.0 1.0 184.5 184.5 184.5 184.5 startRide
1.0s 0 1.0 1.0 0.1 0.1 0.1 0.1 updateActiveRides
2.0s 0 0.5 0.5 13.1 13.1 13.1 13.1 addUser
2.0s 0 0.5 0.5 24.1 24.1 24.1 24.1 addVehicle
2.0s 0 0.5 0.5 75.5 75.5 75.5 75.5 applyPromoCode
2.0s 0 0.5 0.5 14.2 14.2 14.2 14.2 endRide
2.0s 0 55.2 57.4 6.6 14.2 18.9 22.0 readVehicles
2.0s 0 2.0 1.5 46.1 52.4 52.4 52.4 startRide
2.0s 0 6.0 3.5 54.5 113.2 113.2 113.2 updateActiveRides
3.0s 0 0.0 0.3 0.0 0.0 0.0 0.0 addUser
3.0s 0 1.0 0.7 71.3 71.3 71.3 71.3 addVehicle
3.0s 0 0.0 0.3 0.0 0.0 0.0 0.0 applyPromoCode
```



 movr	1.2 MiB	6	39	(n1,n2,n3)
 postgres	0 B	0	0	None
 system	3.3 MiB	48	51	(n1,n2,n3)

```

300.0s      0      17241      57.5      17.4      5.8      56.6      285.2      872.4
PS C:\Users\dangq> cockroach sql --insecure --host=localhost:26257
#
# Welcome to the CockroachDB SQL shell.
# All statements must be terminated by a semicolon.
# To exit, type: \q.
#
# Client version: CockroachDB CCL v23.2.0-beta.1 (x86_64-w64-mingw32, built 2023/11/22 19:28:21
# Server version: CockroachDB CCL v23.1.11 (x86_64-pc-linux-gnu, built 2023/09/27 01:53:43, go1
warning: server version older than client! proceed with caution; some features may not be avail

# Cluster ID: e48d7083-962f-4646-bcd4-866b29d96ee2
#
# Enter \? for a brief introduction.
#
root@localhost:26257/defaultdb> show databases
                                -> show databases;
ERROR: statement ignored: at or near "show": syntax error
SQLSTATE: 42601
DETAIL: source SQL:
show databases
show databases
^
HINT: try \h SHOW DATABASES
root@localhost:26257/defaultdb> show databases;
 database_name | owner | primary_region | secondary_region | regions | survival_goal
-----+-----+-----+-----+-----+-----
 defaultdb    | root  | NULL           | NULL              | {}      | NULL
 movr         | root  | NULL           | NULL              | {}      | NULL
 postgres     | root  | NULL           | NULL              | {}      | NULL
 system       | node  | NULL           | NULL              | {}      | NULL
(4 rows)

Time: 12ms total (execution 8ms / network 3ms)

root@localhost:26257/defaultdb>
M-? toggle key help * C-d erase/stop * C-c clear/cancel * C-r search hist * M-. hide/show prompt

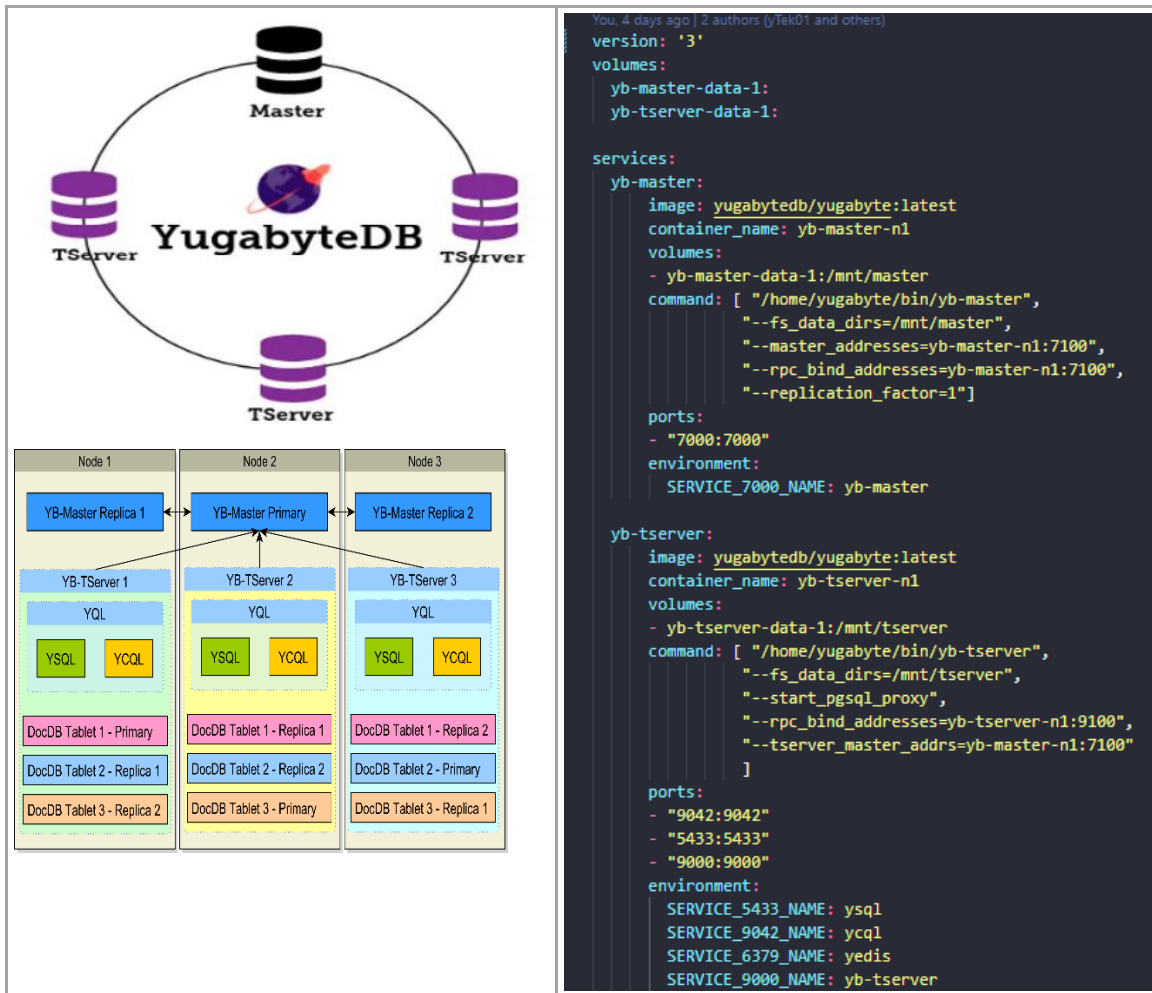
```

YugabyteDB

YugabyteDB is distributed PostgreSQL for enterprise applications delivered as a flexible cloud service.

PostgreSQL Compatibility YugabyteDB is more than just wire compatible with PostgreSQL, it is code compatible. YugabyteDB achieves this by reusing PostgreSQL's query layer to achieve a high degree of compatibility with existing PostgreSQL applications or those that can be migrated to PostgreSQL.

YugabyteDB top-level architecture



Start the Cluster: `docker-compose -f ./docker-compose.Yugabyte.yaml up -d`

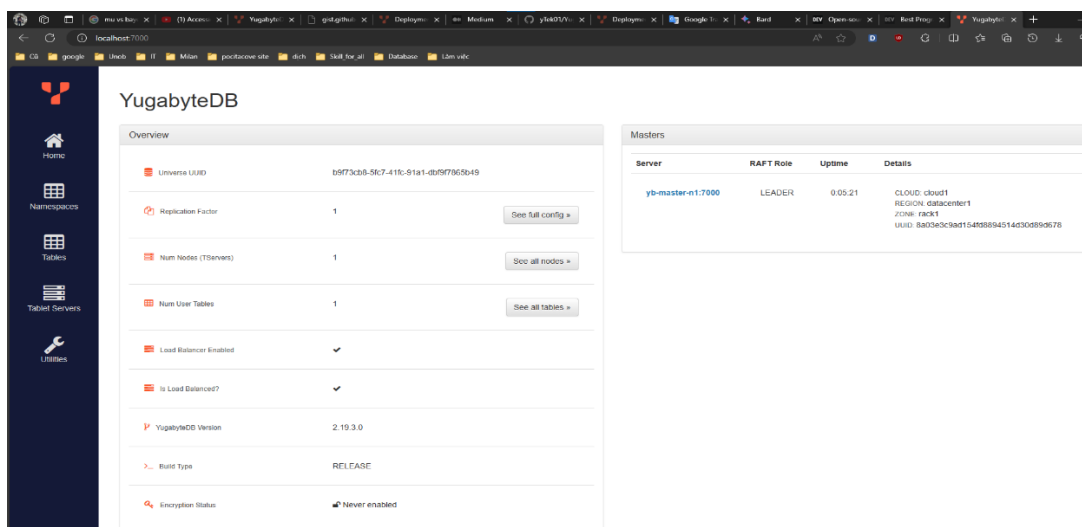
```

[+] Building 0.0s (0/0)
[+] Running 6/6
✓ Network yugabytedb-on-docker_default Created 0.1s
✓ Volume "yugabytedb-on-docker_yb-tserver-data-1" Created 0.0s
✓ Volume "yugabytedb-on-docker_yb-master-data-1" Created 0.0s
✓ Container yb-master-n1 Started 0.5s
✓ Container yb-tserver-n1 Started 0.1s
✓ Container yb-client-n1 Started 0.1s
PS D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker>
  
```

```
See 'docker --help'
PS D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker> docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
c99b043b0fc0        bridge              bridge              local
ea8386050d1c        host                host                local
2b66c271c551        none                null                local
9a2ca7fc5074        roachnet            bridge              local
874129582da1        sql_performance_default bridge              local
4254157d1375        sql_performance_roachnet bridge              local
f73a866a6e4d        tai_cockroach_roachnet bridge              local
a4f4ec1634bc        yugabytedb-on-docker_default bridge              local
PS D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker>
```

```
See 'docker --help'
PS D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker> docker volume ls
DRIVER              VOLUME NAME
local               2d6213d8751ad73d660d2240397130a8b406aa112e4049083f69204b567c82a2
local               2ddd4fa764831b5073226fbc2b55abcfcc78d73e1bdd4f04bcb3c58eba0b04ebc
local               5e978090de4981b348e653f62a0b9bd10df78ebb68bfa1f7d58d03a49cd59399
local               651bbb3883062ac1815ccbaff30b39997774e9855410ed6526a2763009a42c41
local               7785680e60822a9236d75a2c4956e52567e22f769dda98cca558c5c0da1b74b8
local               b190f9a9936f34660de292f465e7c18a991f9a607f34f3d28a0bd054ff681289
local               c701a942c2f3431502a4fb36eea227d7017bc3c36619c8ae4c66d3250ebad0b
local               tai_cockroach_roach1
local               tai_cockroach_roach2
local               tai_cockroach_roach3
local               yugabytedb-on-docker_yb-master-data-1
local               yugabytedb-on-docker_yb-tserver-data-1
PS D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker>
```

Test the APIs: Clients can now connect to the YSQL API at localhost:5433 and the YCQL API at localhost:9042. The yb-master admin service is available at <http://localhost:7000>.



[YugabyteDB metrics](#) | [YugabyteDB Docs](#)

```

4. Connect to YSQL
YCQL and YSQL APIs are enabled by default on the cluster.
```bash
docker exec -it yb-tserver-n1 /home/yugabyte/bin/ysqlsh -h yb-tserver-n1
```

5. Connect to YCQL
```bash
docker exec -it yb-tserver-n1 /home/yugabyte/bin/ycqlsh yb-tserver-n1
```

```

```

PS D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker> docker exec -it yb-tserver-n1 /home/yugabyte/b
ysqlsh (11.2-YB-2.19.3.0-b0)
Type "help" for help.

yugabyte=#

```

```

D:\Documents\Unob_7\STC\STC_code\YugabyteDB-on-Docker>docker exec -it yb-tserver-n1 /home
Connected to local cluster at yb-tserver-n1:9042.
[ycqlsh 5.0.1 | Cassandra 3.9-SNAPSHOT | CQL spec 3.4.2 | Native protocol v4]
Use HELP for help.
ycqlsh>

```

```

9. Load the data into the database.
```bash
\i load_data.sql
```

10. Connect to YSQL, with the new user and database connection.
```bash
docker exec -it yb-tserver-n1 /home/yugabyte/bin/ysqlsh -h yb-tserver-n1 -U postgres -d dvdrental
```

* List all the tables in the dvdrental database.
```bash
\dt
```

* Check the actors table.
```bash
SELECT * FROM public.actors;
```

```